

Academic Talk Summarizer: Speaker Guide

Quick Start

As a speaker, you evaluate how well AI models understood your talk. Your expert evaluation helps improve AI performance on academic content.

Using the Speaker Interface

0. **Login** Use username: speaker, use password speakerpass
1. **Select your talk** from the dropdown
2. **Choose number of questions** to review (default: 5)
3. **Load Questions & Answers** - you'll see audience members asked about your talk
4. **Rate each answer:**
 - **Preferred Answer:** Which model better captured your points? (optional)
 - **Quality Scores:** Rate each model 1-10 based on technical accuracy (optional)
5. **Submit Ratings** - these contribute to the separate "Speaker Leaderboard"

Viewing Results (Leaderboard Tab)

- **Audience Leaderboard:** Shows how general audience rated the models, updated in real time
- **Speaker Leaderboard:** Shows expert evaluations from speakers only, updated in real time

Tips

- Focus on technical accuracy when rating
- Consider whether the AI captured nuanced points from your talk
- Rate based on how well you'd want this summary to represent your work
- Even rating a subset of questions is valuable, they contribute to the separate "Speaker Leaderboard"

Why Speaker Ratings Matter

- You have unique insight into whether LLMs correctly understood your work
- Your ratings help identify which models best handle technical academic content
- The platform tracks speaker vs. audience ratings separately to compare perspectives
- These will be compared with LLM as a judge post-hoc to evaluate the accuracy of this approach in this setting